

# BIG DATA & DISTRIBUTED PROCESSING

- Big Data Overview
- Scalable Systems
- Hadoop
- Spark
  - History
  - RDDs
  - DataFrames
  - Spark Architecture
  - Spark API
- PySpark Examples
- PySpark Exercise
- Assignment


# START SPARK SESSION

```
import pyspark
from pyspark.sql import SparkSession
```

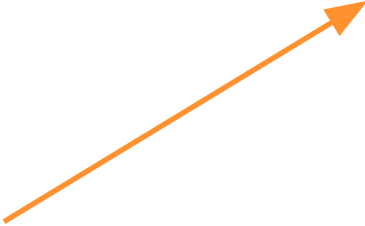
```
conf = pyspark.SparkConf().setAll([
    ('spark.master', 'local[*]'),
    ('spark.app.name', 'PySpark Demo')])
```

```
spark = SparkSession.builder.config(conf=conf).getOrCreate()
```


Use \* to use all available cores, or integer value to specify number of cores to use



Configuration parameters for Spark session



Get existing Spark session or create new one



# LOAD DATA

- Loading data from local file system

```
df = spark.read.text("file:///<path>/<file>.txt")
```

```
df = spark.read.csv("file:///<path>/<file>.csv",  
                    header=True).cache()
```

- Loading data from HDFS

```
df = spark.read.text("hdfs:///<path>/<file>.txt") \  
    .cache()
```

```
df = spark.read.csv("hdfs:///<path>/<file>.csv",  
                    header=True,  
                    interSchema=True).cache()
```

**Indicates whether  
column headers exist**



**Automatically infer  
data types of columns**



**Cache data in memory**



# CHAINING

```
text_file = sc.textFile("hdfs://...")

counts = text_file.flatMap(lambda line: line.split(" ")) \
    .map(lambda word: (word, 1)) \
    .reduceByKey(lambda a, b: a + b)

counts.saveAsTextFile("hdfs://...")
```

**Chaining: Making multiple  
method calls on same object**



**Line continuation  
indicator**



## RDD Wordcount

# CREATE DATAFRAME

```
Employee = Row("name", "dept", "state", "salary")
employee1 = Employee('James', 'Sales', 'CA', 100000)
employee2 = Employee('Mary', 'Finance', 'NY', 120000)
employee3 = Employee('Jane', 'Sales', 'WA', 160000)
employees = [employee1, employee2, employee3]
employeesDF = spark.createDataFrame(employees)
employeesDF.show()
```

name	dept	state	salary
James	Sales	CA	100000
Mary	Finance	NY	120000
Jane	Sales	WA	160000

# SENTENCE DATAFRAME

```
sent_0 = Row(value='This is a sentence')
sent_1 = Row(value='This is another sentence')
sentences = [sent_0, sent_1]
sentenceDF = spark.createDataFrame(sentences)
sentenceDF.show()
```

```
+-----+
|                value|
+-----+
| This is a sentence|
|This is another s...|
+-----+
```

# DATAFRAME OPERATIONS

- Check type - `type(variable)`
- Display schema - [printSchema](#)
- Show content of the DataFrame - [show](#)
- Number of rows - [count](#)
- Number of columns - `len(dataFrame)`
- Select columns - [select](#)
- Summary - [describe](#)
- Group by columns - [groupBy](#)

# DATAFRAME OPERATIONS

- Filter based on condition on columns - [filter](#)
- Sort by column name - [sort](#)
- Split string based on delimiter - [split](#)
- Explode(Map rows to columns) - [explode](#)
- Alias set - [alias](#)



# SPLIT

```
from pyspark.sql.functions import *
```

```
wordsDF1 = sentenceDF.select(split("value", " ").alias("csv"))  
wordsDF1.show()
```

	csv
[This, is, a, sen...	
[This, is, anothe...	

Split each line based on specified delimiter

# EXPLODE

wordsDF1

csv
[This, is, a, sen...]
[This, is, anothe...]

```
wordsDF2 = wordsDF1.select(explode("csv").alias("word"))  
wordsDF2.show()
```

word
This
is
a
sentence
This
is
another
sentence

Map columns to rows

## Combining split() and explode()

```
wordsDF = sentenceDF.select(explode(split("value", " ")).alias("word"))  
wordsDF.show()
```

word
This
is
a
sentence
This
is
another
sentence

# SAVING DATAFRAME TO FILE

DataFrame contents are coalesced into 1 partition and written to `employees_sorted.csv/part-00000-*.csv`

Save column  
headers



Overwrite  
existing file



```
employees_sortedDF.coalesce(1).\n  write.csv("file:///<path>/employees.csv", \n            header=True, mode="overwrite")
```

	name	dept	state	salary
1	Jane	Sales	WA	160000
2	Mary	Finance	NY	120000
3	James	Sales	CA	100000

# SAVING DATAFRAME TO HDFS

DataFrame contents are coalesced into 1 partition and written to HDFS

```
employees_sortedDF.coalesce(1) .\
  write.csv("hdfs:///<path>/employees.csv", \
    header=True, mode="overwrite")
```

	name	dept	state	salary
1	Jane	Sales	WA	160000
2	Mary	Finance	NY	120000
3	James	Sales	CA	100000