

CNN TRANSFER LEARNING EXERCISE

- Data
 - Cats and dogs images from Kaggle
- Exercises
 - Feature extraction
 - ▢ Use pre-trained CNN to extract features from images
 - ▢ Train neural network to classify cats/dogs using extract features
 - Fine tune
 - ▢ Adjust weights of last few layers of pre-trained CNN through training

FEATURE EXTRACTION


- Data

- Cats and dogs images from Kaggle

- Method

- Use VGG16 trained on ImageNet data as pre-trained model. Remove last fully connected layer.
- Extract features from pre-trained model and save
- Neural network then trained on extracted features to classify cats vs. dogs

TRANSFER LEARNING - FEATURE EXTRACTION



Source: <https://www.cs.toronto.edu/~frossard/post/vgg16/>

AWS SETUP

- Create a SageMaker notebook instance
- Notebook instance name
- Notebook instance type(with GPU acceleration)
 - **ml.p2.xlarge**
- IAM Role - Default
- Create notebook instance
 - Click on 'Create notebook instance'
 - Wait until notebook status changes to InService.
 - Click 'Open JupyterLab'
- Upload existing notebook
 - Click on up arrow to upload notebook (features.ipynb and finetune.ipynb)
 - Select **conda_tensorflow2_p36** kernel

LOOK AT DATA

- In terminal window, do the following
- Get counts of images
 - `ls -l ~/train/cats/* | wc -l`
 - `ls -l ~/train/dogs/* | wc -l`
 - `ls -l ~/validation/cats/* | wc -l`
 - `ls -l ~/validation/dogs/* | wc -l`

DATA DESCRIPTION

- Subset of Kaggle cats and dogs dataset
- Train
 - 1000 cats + 1000 dogs
- Validation
 - 200 cats + 200 dogs
- Test
 - 200 cats + 200 dogs



PRINT SOFTWARE VERSIONS

```
import tensorflow as tf  
print (tf.__version__)  
print (keras.__version__)
```



SET DATA PARAMETERS

- Set image dimensions

- o `img_width, img_height = 150, 150`*



- Set data location

- o `train_data_dir = 'train'`*



- o `validation_data_dir = 'validation'`*



- Set number of images

- o `nb_train_samples = 2000`*



- o `nb_validation_samples = 800`*



(150, 150, 3)

METHOD TO EXTRACT FEATURES FROM PRE-TRAINED MODEL

```
def save_features():
```

```
...
```

1. Scale pixel values in each image
2. Load weights for pre-trained network without top classifier
3. Generator reads images from subdir, batch_size number of images at a time.
4. Feed images through pre-trained network and extract features
5. Save features
6. Repeat 3-5 for validation data

CALL METHOD TO EXTRACT & SAVE FEATURES

```
save_features()
```



```
Found 2000 images belonging to 2 classes.
```

```
Found 800 images belonging to 2 classes.
```

Layer (type)	Output Shape	Param #
=====		
input_2 (InputLayer)	(None, None, None, 3)	0
block1_conv1 (Conv2D)	(None, None, None, 64)	1792
block1_conv2 (Conv2D)	(None, None, None, 64)	36928
block1_pool (MaxPooling2D)	(None, None, None, 64)	0
block2_conv1 (Conv2D)	(None, None, None, 128)	73856
block2_conv2 (Conv2D)	(None, None, None, 128)	147584
block2_pool (MaxPooling2D)	(None, None, None, 128)	0

LOAD SAVED FEATURES

- Add name of file containing saved features

- o For train data

train_data = np.load ('features_train.npy')



- o For validation data

validation_data = np.load ('features_validation.npy')



(2000,) (800,)

CREATE TOP MODEL

- Model
 - Fully connected layer from input to hidden
 - 256 nodes in hidden layer
 - Rectified linear activation function
 - Fully connected layer from hidden to output
 - 1 node in output layer (cat or dog)
 - Sigmoid activation function

TRAIN TOP MODEL

- Set number of training iterations
 - epochs = 50 
- Train model, keeping track of history

```
from keras.callbacks import History
hist = top_model.fit(train_data, train_labels,
                    epochs=epochs,
                    batch_size=batch_size,
                    validation_data=(validation_data, validation_labels))
```

```
Train on 2000 samples, validate on 800 samples
Epoch 1/50
2000/2000 [=====] - 1s 451us/step - loss: 0.7173 - acc: 0.7445 - v
al_loss: 0.2955 - val_acc: 0.8788
Epoch 2/50
2000/2000 [=====] - 1s 262us/step - loss: 0.3366 - acc: 0.8525 - v
al_loss: 0.2619 - val_acc: 0.8925
Epoch 3/50
2000/2000 [=====] - 1s 262us/step - loss: 0.3366 - acc: 0.8525 - v
al_loss: 0.2619 - val_acc: 0.8925
```

SAVE MODEL & WEIGHTS

- Add name for model files

0 top_model_file = 'features_model'



- Save model and weights


```
# Save model & weights to HDF5 file
top_model_file = 'features_model'
top_model.save(top_model_file + '.h5')

# Save model to JSON file & weights to HDF5 file
top_model_json = top_model.to_json()
with open(top_model_file + '.json', 'w') as json_file:
    json_file.write(top_model_json)
top_model.save_weights(top_model_file + '-wts.h5')
```

TEST MODEL ON VALIDATION DATA

- Get prediction results on validation data

```
# Results on validation set
print (top_model.metrics_names)
results = top_model.evaluate (validation_data, validation_labels)
print (results)
```




```
['loss', 'acc']
800/800 [=====] - 0s 43us/step
```

- [1.1933523465033795, 0.885]
 - Results should be the same
- Validation accuracy on CNN trained from scratch
 - ~87.5%

TRANSFER LEARNING - FINE TUNING

- Data
 - Cats and dogs images from Kaggle
- Method
 - Use VGG16 trained on ImageNet data as pre-trained model.
 - Replace last fully connected layer with neural network trained from Feature Extraction hands-on.
 - Fine tune last convolution block and fully connected layer.

TRANSFER LEARNING - FINE TUNING



Source: <https://www.cs.toronto.edu/~frossard/post/vgg16/>

SET DATA PARAMETERS

- Set image dimensions

- o `img_width, img_height = 150, 150`*



- Set data location

- o `train_data_dir = 'train'`*



- o `validation_data_dir = 'validation'`*



- Set number of images

- o `nb_train_samples = 2000`*



- o `nb_validation_samples = 800`*



(150, 150, 3)

LOAD PRE-TRAINED MODEL


- Load pre-trained model without last fully connected layer

```
base_model = applications.VGG16  
(weights='imagenet',  
include_top=False,  
input_shape=(img_width,img_height,3))  
print ('Model loaded')
```

- Print out base model summary

```
base_model.summary()      
```

TRANSFER LEARNING - FINE TUNING



Source: <https://www.cs.toronto.edu/~frossard/post/vgg16/>

CREATE TOP MODEL

- Create top model
 - Create fully connected layer as top model and connect to pre-trained base model
- Load top model's weights
 - Weights are in 'features_model_wts.h5'
- Add top model to base CNN to create model
- Freeze weights

for layer in model.layers[:15]

layer.trainable = False



- Compile model
 - Print out model summary
- model.summary()*



MODEL

- Original Model

Total params: 14,714,688

Trainable params: 14,714,688

Non-trainable params: 0

- Freeze some weights

```
# Freeze weights in CNN up to last Conv block
for layer in model.layers[:15]:
    layer.trainable = False
```

Trainable params: 9,177,089

Non-trainable params: 7,635,264

PREPARE DATA

- Set batch size

batch_size = 16



- Set batch size for `train_generator`

```
train_generator = train_datagen.flow_from_directory(
    train_data_dir,
    target_size=(img_width, img_height),
    batch_size=batch_size,
    class_mode='binary',
    seed=seed)
```



FINE TUNING

- Set number of training epochs

epochs = 5



- Set batch size for train_generator

from keras.callbacks import History

hist = model.fit_generator(

train_generator,

steps_per_epoch = nb_train_samples // batch_size,

epochs = epochs,



validation_data = validation_generator,


validation_steps = nb_validation_samples // batch_size,


initial_epoch=0,

verbose = 2)

GET RESULTS

- Get classification results after fine tuning

```
results = model.evaluate_generator(  
    train_generator,   
    steps=nb_train_samples // batch_size)  
print (results)
```

```
results = model.evaluate_generator(  
    validation_generator,   
    steps=nb_validation_samples // batch_size)  
print (results)
```

SAVE MODEL & WEIGHTS

- Save model & weights

```
model_file = 'finetune'
```



- Get results on validation set

```
print (model.metrics_names)
```

```
results = model.evaluate_generator(
```



```
validation_generator,
```

```
steps = nb_validation_samples // batch_size)
```

```
print (results)
```

OUTPUT TRAINING HISTORY

- Print history

print (hist.history)



INFERENCE

- Use model to predict class of image

```
result = model.predict(x)
```



```
print ("Prediction probability: ", result)
```



Fine Tuning Results

- Before fine tuning
 - [loss, accuracy]:
[0.641176361694085, 0.925]
[1.1933523442077918, 0.885]
- After fine tuning
 - Train adjustable parameters for 5 epochs
 - [loss, accuracy]:
[0.07257955298712478, 0.978]
[0.29664344725897535, 0.9125]