

AWS SageMaker STEPS

- Create S3 bucket
 - Upload data to bucket
- Create SageMaker notebook instance
- Create Jupyter notebook
 - Set up data
 - Create training job to train model
 - Fit model to data
 - Deploy model for inference
 - Perform inference using deployed model
 - Evaluate model
- Clean up

CREATE BUCKET

- New bucket
 - Do this if you want to create a new bucket
- Services -> S3
- Create bucket
 - Enter bucket name
 - Bucket name must be unique across all S3 buckets
 - Cannot contain uppercase characters or underscores
 - Must start with lowercase letter or number
 - More on bucket naming
 - ❖ <https://docs.aws.amazon.com/AmazonS3/latest/dev//BucketRestrictions.html#bucketnamingrules>
 - Select US West for region (default)
 - Click Create

CREATE FOLDER

- Existing bucket
 - If using an existing bucket, then just create folder within bucket.
- Select bucket
 - Services -> S3
 - Click on bucket to use
- Create folder
 - Select bucket
 - Select 'Create folder' for bucket
 - Enter name for folder
 - Click 'Save'

UPLOAD DATA TO S3 BUCKET

- Services -> S3
- Create folder (if needed)
- Click into folder
- Click 'Upload'
- Select files
 - Drag & drop or select files to upload
 - Click 'Upload'

AMAZON SAGEMAKER NOTEBOOK INSTANCE

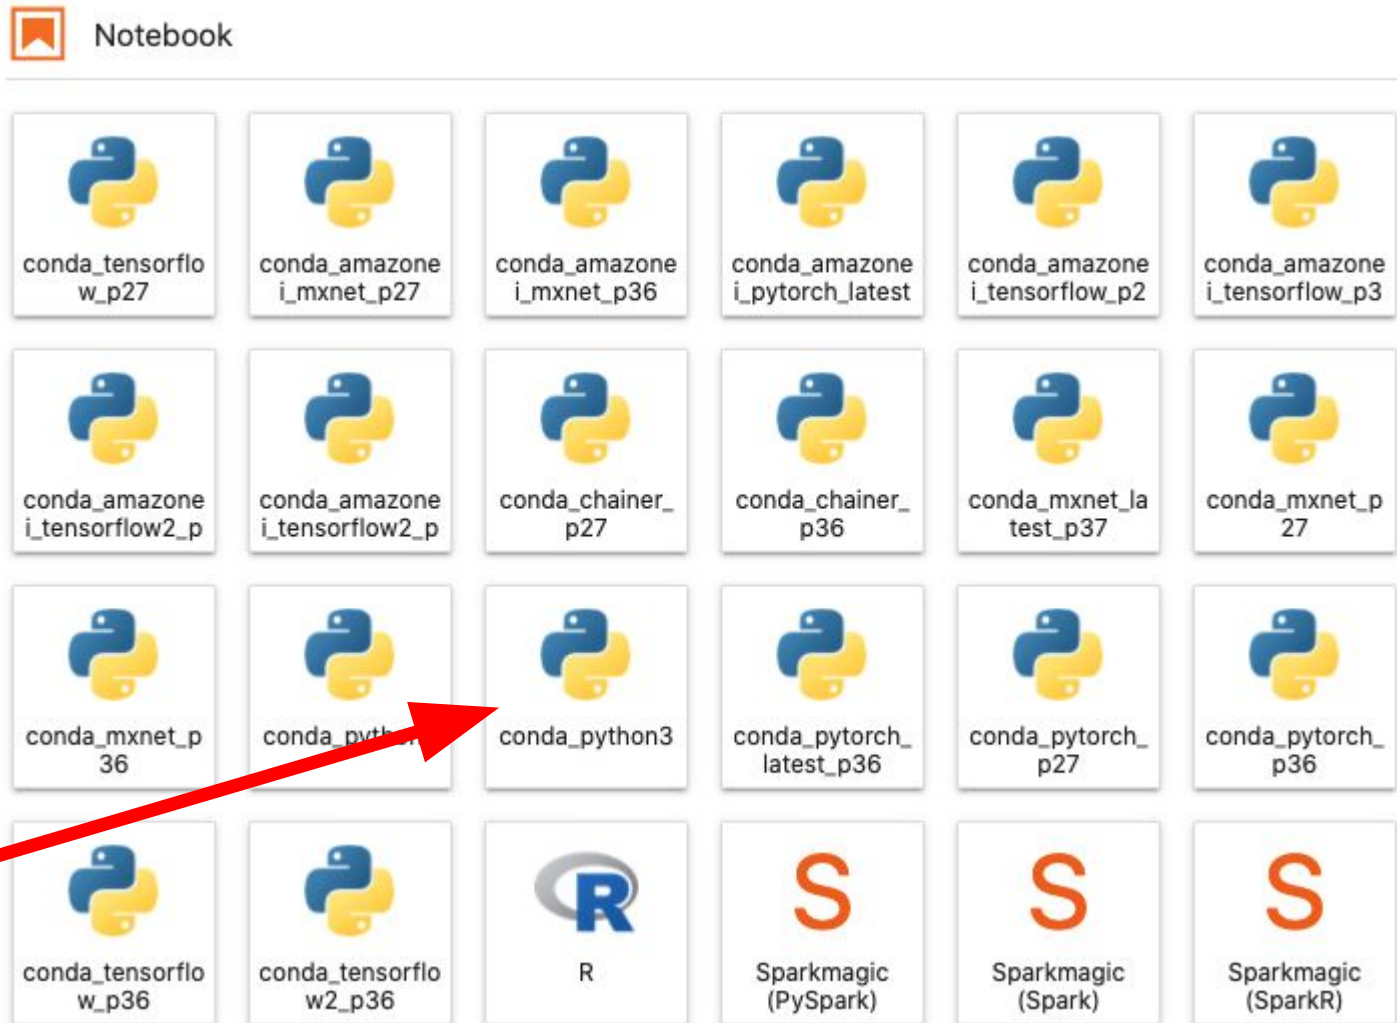
- Services -> Amazon SageMaker (under Machine Learning)
- Click on Dashboard in the left menu
- Click on Notebook Instances
- Click on 'Create notebook instance'
- Notebook instance name
 - Enter notebook instance name
- Notebook instance type
 - **ml.t3.medium**
- IAM Role
 - Select 'Create a new role'
 - Specific S3 bucket → Enter your bucket name
 - Click 'Create role'

CREATE SageMaker NOTEBOOK INSTANCE

- Create notebook instance
 - Click on 'Create notebook instance'
 - Wait until notebook status changes to InService.
 - May take a few minutes. Click on arrow to refresh.
 - Click 'Open JupyterLab'
 - Upload existing notebook
 - Click on up arrow to upload notebook
 - Select **conda_python3** kernel
 - Create new notebook
 - File -> New -> Notebook
 - For 'Select Kernel', choose conda_python3
 - Edit notebook

CREATE JUPYTER NOTEBOOK

- Starting new notebook from within JupyterLab
- For this class, use conda-python3 as kernel



SPECIFY DATA LOCATION ON S3

```
bucket = "<bucket_name>"
```

```
prefix = "<folder_name>"
```

```
fname = "<file_name>"
```

```
data_fname = "s3://{}/{}/{}/{}".format(bucket,prefix,fname)
```

```
df = pd.read_csv(data_fname)
```


SET UP DATA AND OUTPUT ON S3 FOR MODEL

```
bucket = "<bucket_name>"
```

```
prefix = "<folder_name>"
```

```
train_path =
```

```
    "{}/{}/tmodel_data/{}".format(bucket, prefix, "train_data.csv")
```

```
train_df.to_csv(train_path, index=False, header=False)
```

```
s3_input_train =
```

```
    sagemaker.inputs.TrainingInput(s3_data=train_path,  
                                   content_type='csv')
```

```
output_location = "s3:///{}/{}/model".format(bucket, prefix)
```

TRAIN MODEL

```
sess = sagemaker.Session()
```

```
from sagemaker.amazon.amazon_estimator import image_uris
xgb_image = image_uris.retrieve(framework="xgboost",
                                region=my_region, version='latest')
```

```
xgb_model = sagemaker.estimator.Estimator(
    xgb_image, iam_role,
    train_instance_count=1,
    train_instance_type='ml.m5.xlarge',
    output_path=<output_location>,
    sagemaker_session=sess)
xgb_model.fit({'train': <s3_train_loc>, 'validation': <s3_val_loc>})
```

REAL-TIME INFERENCE

```
xgb_predictor = xgb_model.deploy (  
    initial_instance_count=1,  
    serializer = sagemaker.serializers.CSVSerializer(),  
    instance_type='ml.t2.medium')
```

BATCH INFERENCE

```
xgb_transformer =  
    xgb_model.transformer(  
        instance_count=1,  
        instance_type='ml.m5.large',  
        output_path=<output_location>)
```

HYPERPARAMETER TUNING

Specify tuning job parameters

```
hyperparameter_ranges = {  
    'max_depth': IntegerParameter(1, 10)}
```

Create tuning job

```
Optimizer = sagemaker.tuner.HyperparameterTuner(  
    estimator=xgb_model,  
    hyperparameter_ranges=hyperparameter_ranges,  
    base_tuning_job_name='XGBoost-Tuner',  
    objective_type='Minimize',  
    # objective_metric_name='validation:accuracy',  
    objective_metric_name='validation:merror',  
    max_jobs=10,  
    max_parallel_jobs=5)
```

Launch tuning job

```
Optimizer.fit({'train': s3_input_train, 'validation': s3_input_val})
```

CLEANING UP AFTER SESSION

- At end of each session
- SageMaker Dashboard
 - Stop notebook instances
 - Delete training jobs
 - Delete hyperparameter tuning jobs
 - Delete inference models
 - Delete batch transform jobs
 - i.e., nothing in green
- S3
 - Download results file if haven't done so

WORK ON EXISTING NOTEBOOK

- Go to SageMaker Dashboard
 - Services -> Amazon SageMaker (under Machine Learning)
- Click on 'Notebook instances'
- Select existing notebook
 - Click on 'Start' under Actions
- Wait until Status changes to InService
- Click on 'Open in JupyterLab'
 - Click on circular arrow to refresh

CLEANING UP AFTER ASSIGNMENT

- SageMaker
 - Delete notebook instances, models, endpoints
- S3
 - Download results file if haven't done so

CLEANING UP AT END OF COURSE

- Delete everything when done
 - Any running service, even if idle, will incur charges!
- SageMaker
 - Delete notebook instances, models, endpoints
- S3
 - Download results file if haven't done so
 - Select folder and select Actions -> Delete
 - If done, also delete S3 bucket
 - **Note: All contents will be lost!**
- All services
 - Delete all other services